

Comprehensive Architectural Framework for Implementing Generative UI in Flutter Document Signing Applications

1. Executive Summary

The transition from imperative, static user interfaces to Generative UI (GenUI) represents a fundamental shift in software engineering, particularly for complex, state-dependent workflows such as digital document signing. The user's request involves a sophisticated requirement: analyzing the architectural patterns presented in modern Generative UI tutorials (specifically leveraging the Vercel AI SDK paradigm demonstrated in the reference video) and transposing these concepts into a production-grade Flutter application. This application requires distinct handling of three specific business flows: **Single Sign (Self-Sign)**, **1-on-1 Sign**, and **Multi-Party Sign**, each characterized by unique constraints regarding recipient cardinality and workflow progression.

This report provides an exhaustive, step-by-step analysis and implementation guide. It synthesizes the theoretical foundations of Generative UI—where Artificial Intelligence (AI) acts as a runtime orchestrator rather than a passive data processor—with the practical realities of the Flutter genui ecosystem. By analyzing the provided research materials, we establish that while the reference video utilizes React and Next.js, the underlying architectural principles—**Tool Calling**, **Zod/JSON Schema validation**, **Server-Driven UI**, and **Streaming State Management**—are directly applicable to Flutter through the genui, json_schema_builder, and Generative_ai packages.¹

The following sections detail the robust engineering of a "Catalog-Based" architecture where the AI agent, guided by rigorous system instructions, dynamically assembles the signing interface. This approach allows the application to enforce strict business rules (e.g., preventing a third recipient in a 1-on-1 flow) not through rigid if-else chains, but through intelligent, context-aware state validation.

2. Architectural Analysis of the Reference Paradigm

2.1 Deconstructing the Video: The Vercel AI SDK Pattern

The reference video (<https://www.youtube.com/watch?v=K2p5Nrn2OSU>) creates a foundation for understanding Generative UI by demonstrating the Vercel AI SDK. Although the video

focuses on web technologies (React/Next.js), the core architectural pattern it advocates is platform-agnostic and critical for our Flutter implementation. The video illustrates a system where the UI is not pre-rendered but is "streamed" as a reaction to AI tool calls.¹

In the video's model, the interaction loop functions as follows:

1. **Intent Recognition:** The user expresses a need (e.g., "I need to check my wallet balance").
2. **Server Action/Tool Call:** The AI identifies a relevant tool (get_balance) and executes it.
3. **Component Generation:** Instead of returning raw text JSON, the server generates a specific React component (e.g., a Pie Chart) populated with that data.²
4. **Client Rendering:** The client streams this component code (or a descriptor) and renders it instantly.

Implications for Flutter:

To replicate this "Right Way" in Flutter, we must substitute the React-specific components with Flutter equivalents. Flutter does not "stream components" in the same way React Server Components (RSC) do. Instead, the Flutter GenUI architecture relies on a Catalog of Widgets. The AI sends a JSON descriptor (using the A2UI protocol), and the Flutter client acts as a factory, matching the descriptor to a pre-built widget in the local catalog.³ This distinction is vital: we are not generating Dart code on the fly; we are generating configurations for existing, high-quality, brand-compliant widgets.⁶

2.2 The Flutter GenUI Ecosystem

The genui package for Flutter implements this exact orchestration layer. It coordinates the flow between the user, the widget tree, and the AI agent. The architecture consists of four pillars that will form the backbone of the signing app⁷:

Component	Function in Signing App	Flutter Implementation
GenUiConversation	The "Brain." Manages history and state.	GenUiConversation class
ContentGenerator	The "Adapter." Connects to LLM (Gemini/Vertex).	FirebaseAiContentGenerator or A2uiContentGenerator
Catalog	The "Toolbox." List of available UI widgets.	Catalog containing CatalogItems
DataModel	The "Memory." Stores session state (URL,	DataModel (Observable JSON store)

	recipients).	
--	--------------	--

The "Right Way" to implement this involves ensuring that the AI never directly manipulates the UI pixels but instead manipulates the DataModel and requests widgets that *bind* to that model. This separation of concerns ensures that the AI cannot "hallucinate" a UI that breaks the app, as it is constrained to the strictly defined Catalog.⁹

3. Designing the Signing Application Architecture

3.1 Workflow Definition

The application must support three distinct flows. In a traditional app, these would be three separate Navigator routes. In GenUI, they are a single "Surface" that adapts based on the DataModel.

1. Self Sign (Single Sign):

- *Constraint*: Recipient List Length = 1.
- *Sequence*: Upload → Review → Sign.
- *Role*: The user is both Sender and Signer.

2. 1-on-1 Sign:

- *Constraint*: Recipient List Length = 2 (Sender + 1 Other, or 2 Others).
- *Sequence*: Upload → Add Recipient (x1) → Review → Send.

3. Multi-Party Sign:

- *Constraint*: Recipient List Length \geq 3.
- *Sequence*: Upload → Add Recipients (xN) → Review → Send.

3.2 The Data Model Schema

The central nervous system of our GenUI implementation is the JSON schema of the DataModel. This schema dictates what the AI "knows" about the current state.

JSON

```
{
  "session": {
    "id": "uuid-v4",
    "status": "initializing | collecting_data | reviewing | sent",
    "flowType": "self_sign | 1_on_1 | multi_party",
    "document": {
      "id": "document-1"
    }
  }
}
```

```

    "url": "https://storage.googleapis.com/...",
    "fileName": "contract.pdf",
    "mimeType": "application/pdf"
},
"recipients": [
{
  "name": "Alice",
  "email": "alice@example.com",
  "role": "signer",
  "order": 1
},
],
"constraints": {
  "minRecipients": 1,
  "maxRecipients": 1
}
}
}

```

Insight: By including a constraints object in the data model, we allow the AI to dynamically adjust the validation rules. When the user selects "Self Sign," the AI updates maxRecipients to 1. The RecipientForm widget then binds to this value, automatically disabling the "Add" button when the limit is reached, providing instant client-side feedback without a server round-trip.⁴

4. Step-by-Step Implementation Strategy

Phase 1: Environment Configuration and Dependencies

To build this the "right way" for a production environment, we must use the `genui_firebase_ai` package. This allows us to use Firebase for secure API key management and authentication, which is critical for a document signing app handling sensitive legal documents.⁴

Step 1.1: pubspec.yaml Configuration

We must include the core GenUI packages and the schema builder. Note that because GenUI is experimental, we may need to reference specific git repositories or alpha versions.³

YAML

```
dependencies:
```

```
  flutter:
```

```

sdk: flutter
# Core GenUI
genui: ^0.1.0-alpha
genui_firebase_ai: ^0.1.0-alpha

# Schema Definition
json_schema_builder:
  git:
    url: https://github.com/flutter/genui.git
    path: packages/json_schema_builder

# Infrastructure
firebase_core: ^3.0.0
firebase_storage: ^12.0.0
file_picker: ^8.0.0
provider: ^6.1.1

```

Step 1.2: Initialization

Initialize Firebase and the GenUiConversation in your main provider. The GenUiConversation is the entry point that encapsulates the connection to the LLM.⁷

Phase 2: Building the Component Catalog

The Catalog is the vocabulary we provide to the AI. We need to define specific widgets for our signing flows. Each widget requires a Schema (input definition) and a WidgetBuilder (rendering logic).⁸

Widget 1: The Flow Selector

This widget allows the user to determine the context. It binds to the flowType field in the DataModel.

- **Schema:**
 - title: String (Label for the card)
 - options: Array of Strings (The available flows)
 - targetPath: String (Where to write the selection in the DataModel)
- **Implementation logic:** When the user selects "1-on-1", the widget updates session.flowType to 1_on_1. Crucially, this update is sent back to the AI as a DataModelUpdate event, triggering the AI to calculate the next step (which would be requesting the document upload).⁴

Widget 2: The Document Uploader (Critical Implementation Detail)

The video analysis suggests using "Server Actions" for complex tasks. In Flutter GenUI, file uploading requires careful handling because LLMs cannot process raw PDF bytes efficiently.

The Right Way:

1. **Client-Side Upload:** The DocumentUploader widget uses file_picker to select the file.
 2. **Direct-to-Storage:** The widget uploads the file directly to Firebase Storage/AWS S3.
 3. **URL Binding:** Upon success, the widget writes the *download URL* and *metadata* (name, size) to the DataModel at session.document.¹¹
 4. **AI Notification:** The DataModel update notifies the AI that document.url is no longer null.
- **Schema:**
 - allowedExtensions: Array(['pdf', 'docx'])
 - storagePath: String (Bucket path)
 - outputUrlPath: String (DataModel path for the result)

Widget 3: The Recipient Manager

This is the most complex widget as it needs to adapt to the three flows.

- **Schema:**
 - recipientsListPath: String (Path to array)
 - minCount: Integer (Bound to session.constraints.minRecipients)
 - maxCount: Integer (Bound to session.constraints.maxRecipients)
- **Dynamic Behavior:**
 - If maxCount == 1 (Self Sign), the "Add Recipient" button disappears after one entry.
 - If minCount == 3 (Multi-Party), the "Continue" button is disabled until 3 items exist.
 - This validation happens *locally* in the widget based on the schema, providing immediate UI feedback.⁴

Phase 3: The Agentic Brain (System Instructions)

The ContentGenerator requires a System Instruction. This is where we implement the business logic for the three flows. Instead of writing Dart code to manage the flow, we write natural language rules.⁴

Drafting the System Prompt:

"You are a Document Signing Orchestration Agent. Your goal is to guide the user through a signing ceremony. You have access to a DataModel containing a 'session' object.

Protocol:

1. **Initialization:** Check if session.flowType is set. If not, generate a FlowSelector widget.
2. **Configuration:**
 - If flowType is 'self_sign', set constraints.maxRecipients = 1.
 - If flowType is '1_on_1', set constraints.maxRecipients = 2 and minRecipients = 2.

- If flowType is 'multi_party', set constraints.minRecipients = 3.
- 3. **Document Acquisition:** If session.document.url is empty, generate a DocumentUploader.
- 4. **Recipient Collection:** If document is present, generate a RecipientManager. configure it using the constraints from step 2.
- 5. **Review:** Once the recipient list satisfies the minRecipients constraint, generate a ReviewCard.
- 6. **Finalization:** If the user approves the review, call the generateSigningUrl tool."

This prompt ensures that the AI strictly adheres to the definitions of the flows provided in your query. It actively monitors the state and only advances when criteria are met.

Phase 4: Integration and The "Send URL" Tool

The final requirement is to "send a url for the other recipients to sign." In GenUI terms, this is a **Tool Call** (or Function Call), not a UI widget.⁶

Defining the Tool:

We define a Dart function generateAndSendEnvelope.

Dart

```
Future<Map<String, dynamic>> generateAndSendEnvelope(Map<String, dynamic> sessionData)
async {
  // 1. Call backend API (e.g., DocuSign API or custom backend)
  final envelopId = await api.createEnvelope(sessionData['document']['url']);

  // 2. Generate signing links
  final links = await api.getRecipientViews(envelopId, sessionData['recipients']);

  // 3. Return results
  return {'status': 'sent', 'signingUrls': links};
}
```

We register this tool with the ContentGenerator. When the user clicks "Send" on the ReviewCard, the widget triggers a UserActionEvent. The AI receives this, recognizes the intent to finalize, and invokes the generateAndSendEnvelope tool. The result is then used to generate a final SuccessCard showing the URLs.⁸

5. Detailed Logic for Specific Flows

5.1 Self-Sign Flow Implementation

- **User Selection:** User taps "Self Sign".
- **AI Logic:**
 - Updates constraints to {min: 1, max: 1}.
 - Renders RecipientManager.
- **User Action:** Adds themselves.
- **Constraint Check:** The list length is 1. max is 1. The "Add" button is disabled. The "Next" button is enabled.
- **Transition:** User taps "Next". AI generates ReviewCard.
- **Completion:** User taps "Sign Now". AI calls tool selfSignDocument().

5.2 1-on-1 Flow Implementation

- **User Selection:** User taps "1-on-1".
- **AI Logic:**
 - Updates constraints to {min: 2, max: 2}.
- **User Action:** Adds Recipient A.
- **Validation State:** List length is 1. min is 2. The RecipientManager widget (client-side) keeps the "Next" button *disabled* and perhaps shows a helper text: "Please add 1 more signer."
- **User Action:** Adds Recipient B.
- **Validation State:** List length is 2. Button enables.
- **Transition:** AI observes valid state and moves to Review.

5.3 Multi-Party Flow Implementation

- **User Selection:** User taps "Multi-Party".
- **AI Logic:**
 - Updates constraints to {min: 3, max: 99}.
- **Behavior:** The RecipientManager enforces that at least 3 people are added.
- **Scaling UI:** For this flow, the AI might choose to parameterize the RecipientManager to show a "Bulk Upload CSV" button (if available in the catalog), recognizing that adding 10 people manually is tedious. This demonstrates the power of GenUI: adapting the interface complexity to the task.⁶

6. Technical Deep Dive: The A2UI Protocol & State Management

6.1 The A2UI Protocol

Understanding the Agent-to-UI (A2UI) protocol is essential for debugging. The messages flowing between your ContentGenerator and the UI are JSON-based instructions.¹⁵

Example A2UI Message for Rendering:

JSON

```
{  
  "type": "render_surface",  
  "surfaceId": "main_flow",  
  "components":  
}
```

6.2 Data Binding and Reactivity

The genui package uses a reactive binding system.

- **subscribeToString:** Widgets use this to listen to specific paths in the DataModel.⁹
- **Optimization:** When session.recipients changes, only the RecipientManager and the AI agent are notified. The header, footer, or other unrelated widgets do not rebuild. This ensures high performance even with complex document structures.

7. Advanced Considerations and Best Practices

7.1 Security and PII

In a signing app, you are handling names, emails, and contracts.

- **Warning:** Never send the full PDF content to the LLM context window. It is slow, expensive, and a privacy risk.
- **Solution:** Only send the *URL* and *metadata* (page count, signature fields count) to the DataModel. The AI reasons about the *metadata*, not the content.
- **Firebase App Check:** Ensure your genui_firebase_ai integration uses App Check to prevent unauthorized access to your billing quota.⁴

7.2 Handling "Hallucinations"

If the AI generates a widget that doesn't exist (e.g., SuperSignatureWidget), the genui package will throw a schema validation error.

- **Mitigation:** Use strict TypeScript-like definitions in your json_schema_builder.

- **Recovery:** Implement an error boundary in Flutter. If a GenUI surface crashes, catch the error and send a hidden message to the AI: "*The previous widget failed to render. Please try again using only standard catalog items.*".⁹

7.3 Persistence

The GenUiConversation state is typically ephemeral. For a signing app, users might leave and return.

- **Implementation:** You must serialize the DataModel to local storage (using shared_preferences or hive) whenever it changes.
 - **Restoration:** On app startup, read the JSON and inject it back into the GenUiManager's initial state. This allows the conversation to resume exactly where it left off.¹⁷
-

8. Conclusion

Implementing Generative UI for a document signing application transforms the user experience from a rigid wizard to a fluid, intelligent dialogue. By leveraging the Flutter genui package, we move the complexity of the three signing flows (Single, 1-on-1, Multi) out of the hard-coded widget tree and into the System Instructions.

The "Right Way" to implement this involves:

1. **Defining a robust DataModel** that acts as the single source of truth for the session state.
2. **Building a flexible Catalog** of atomic widgets (Uploader, Form, List) that enforce constraints passed via properties.
3. **Crafting strict System Instructions** that map the business rules of the three flows to specific constraints in the DataModel.
4. **Isolating File I/O** to client-side widgets, sharing only URLs with the AI agent to maintain performance and privacy.

This architecture ensures your application is scalable, maintainable, and capable of adapting to future requirements (e.g., adding a "Witness Sign" flow) simply by updating the AI's instructions rather than refactoring the codebase.

Works cited

1. A Complete Guide to Vercel's AI SDK - Codecademy, accessed December 25, 2025, <https://www.codecademy.com/article/guide-to-vercel-s-ai-sdk>
2. Vercel AI SDK Overview Tutorial | Generative UI, Streaming, Agentic Tool Functions, accessed December 25, 2025, <https://www.youtube.com/watch?v=D48I3Nd0E5U>
3. flutter/genui - GitHub, accessed December 25, 2025, <https://github.com/flutter/genui>

4. Get started with the GenUI SDK for Flutter, accessed December 25, 2025,
<https://docs.flutter.dev/ai/genui/get-started>
5. Multi-Step & Generative UI | Vercel Academy, accessed December 25, 2025,
<https://vercel.com/academy/ai-sdk/multi-step-and-generative-ui>
6. GenUI SDK for Flutter, accessed December 25, 2025,
<https://docs.flutter.dev/ai/genui>
7. GenUI SDK main components and concepts - Flutter documentation, accessed December 25, 2025, <https://docs.flutter.dev/ai/genui/components>
8. genui | Flutter package - Pub.dev, accessed December 25, 2025,
<https://pub.dev/packages/genui>
9. How to Use GenUI in Flutter to Build Dynamic, AI-Driven Interfaces - freeCodeCamp, accessed December 25, 2025,
<https://www.freecodecamp.org/news/how-to-use-genui-in-flutter-to-build-dynamic-ai-driven-interfaces/>
10. README.md - flutter/genui - GitHub, accessed December 25, 2025,
<https://github.com/flutter/genui/blob/main/packages/genui/README.md>
11. flutter_file_uploader - Dart API docs - Pub.dev, accessed December 25, 2025,
https://pub.dev/documentation/flutter_file_uploader/latest/
12. Upload files - Flutter - AWS Amplify Gen 1 Documentation, accessed December 25, 2025, <https://docs.amplify.aws/gen1/flutter/build-a-backend/storage/upload/>
13. Build a form with validation - Flutter documentation, accessed December 25, 2025, <https://docs.flutter.dev/cookbook/forms/validation>
14. Tool calls (aka function calls) - Flutter documentation, accessed December 25, 2025, <https://docs.flutter.dev/ai-best-practices/tool-calls-aka-function-calls>
15. Introducing A2UI: An open project for agent-driven interfaces - Google for Developers Blog, accessed December 25, 2025,
<https://developers.googleblog.com/introducing-a2ui-an-open-project-for-agent-driven-interfaces/>
16. Input and events - Flutter documentation, accessed December 25, 2025,
<https://docs.flutter.dev/ai/genui/input-events>
17. Read and write files - Flutter documentation, accessed December 25, 2025,
<https://docs.flutter.dev/cookbook/persistence/reading-writing-files>